

Analysis of the television data set

Arthur Allignol

We read the data using the following command

```
tele <- read.csv("/data/Ulm/Teaching/SS_2014/Intro_Biostat_Comp/Illustrations/Television/data/television.csv",
                 stringsAsFactors = FALSE, na.strings = "*")
head(tele)
```

##	Country	life	tv	phys	fem	male
## 1	Argentina	70.5	4.0	370	74	67
## 2	Bangladesh	53.5	315.0	6166	53	54
## 3	Brazil	65.0	4.0	684	68	62
## 4	Canada	76.5	1.7	449	80	73
## 5	China	70.0	8.0	643	72	68
## 6	Colombia	71.0	5.6	1551	74	68

The aim of the analysis is to assess whether there exists a relation between the number of people per television set and life expectancy.

We first note a couple of facts.

- Japan is the country with the highest life expectancy while Ethiopia has the lowest;
- “Unsurprisingly”, the United States have the lower number of habitants per television set and is equal to 1.3.

We will first explore the data graphically in Section 1 and perform a simple analysis in Section 2.

1 Graphical exploration

We first display a scatterplot of `life` and `tv` in Figure 1). A trend is clearly to be seen but is non-linear.

```
ggplot(tele, aes(x = tv, y = life)) +
  geom_jitter(size = 3) + geom_smooth(method = "lm") +
  theme_economist()

## Warning: Removed 2 rows containing missing values (stat_smooth).
## Warning: Removed 2 rows containing missing values (geom_point).
```




Figure 1: Scatterplot of life expectancy and number of individuals per television set

As a second step, we log-transform the variable `tv`. The relation between life expectancy and log televisions is in Figure 2. The transformation leads to an (almost) perfect linear relation between these two variables.

```
tele$log_tv <- log(tele$tv)
ggplot(tele, aes(x = log_tv, y = life)) +
  geom_jitter(size = 3) + geom_smooth(method = "lm") +
  theme_economist()

## Warning: Removed 2 rows containing missing values (stat_smooth).
## Warning: Removed 2 rows containing missing values (geom_point).
```

2 Linear model

We fit a linear model with life expectancy as outcome and the log-transformed number of individuals per television set.




Figure 2: Scatterplot of life expectancy and log- number of individuals per television set

```
fit_lm <- lm(life ~ log(tv), tele)
```

The results are displayed in table 1.

```
out <- summary(fit_lm)$coefficients
out[, 1:3] <- round(out[, 1:3], 2)
out[, 4] <- format.pval(out[, 4], digits = 2,
                        eps = 10-3)

print(xtable(out, caption = "Linear model",
             label = "tab:tab_lm"),
      caption.placement = "top")
```

We can check the model fit by plotting the residuals versus the fitted values (Figure 3).

```
df <- fortify(fit_lm)
ggplot(df, aes(.fitted, .resid)) +
  geom_point(size = 2) +
  geom_smooth(se=FALSE) +
```

Table 1: Linear model

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	77.89	1.22	63.83	<0.001
log(tv)	-4.26	0.43	-9.9	<0.001

```

scale_x_continuous("Fitted Values") +
scale_y_continuous("Residual") +
labs(title = "Residuals vs fitted") +
theme_economist()

## geom_smooth: method="auto" and size of largest group is <1000, so using loess.
Use 'method = x' to change the smoothing method.

```

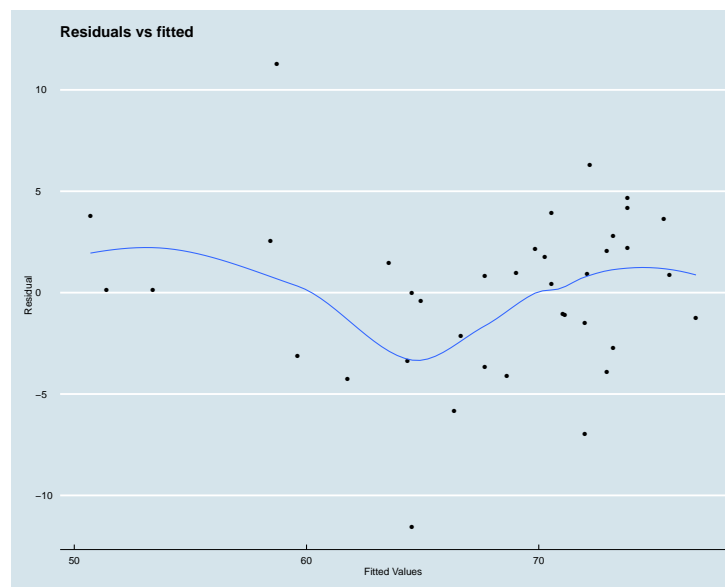


Figure 3: Residuals VS fitted